



Distance-based Classifiers

Department of Computer Languages and Systems

Introduction: general context


Non-parametric, supervised models can be applied if there exist ...

- *objects* capable of being described by some formal representation (a vector, a string, ...)
- *classes* or categories that group the objects
- need to *automate the classification* (or identification) of new objects into some known class

Introduction: an example of classification

A real problem: identification of oil-spills in the ocean


- **objects**: satellite images
- **formal representation**: vector of features or attributes (image *pixels*, *image features*, etc.)
- **classes**: oil-spills, look-alikes (ships, islands, etc.)



Distance-based classifiers

Motivation:

- the best model is not always the most complicated
- these classifiers are based on the idea that **objects of the same class are the most similar**
- they use some **similarity (distance) measures** defined on attribute vectors



Minimum distance classifier

This is by far the simplest classification model

- it does **not take all the training data into account**
- it measures the similarity between test data and some **representative prototype** of each class
- a common **prototype is the centroid** of each class
- some distance is calculated for each prototype and the **shortest distance** is determined
- a test object is assigned to the **class of the nearest centroid** (the minimum distance)

Minimum distance classifier (ii)

minimum_distance(T_{tra} , x)

min \leftarrow MAX

for each class ω_i

$z_i \leftarrow \text{compute_centroid}(\omega_i)$


for each prototype z_i

dist $\leftarrow \text{compute_distance}(x, z_i)$

if dist < min

min \leftarrow dist

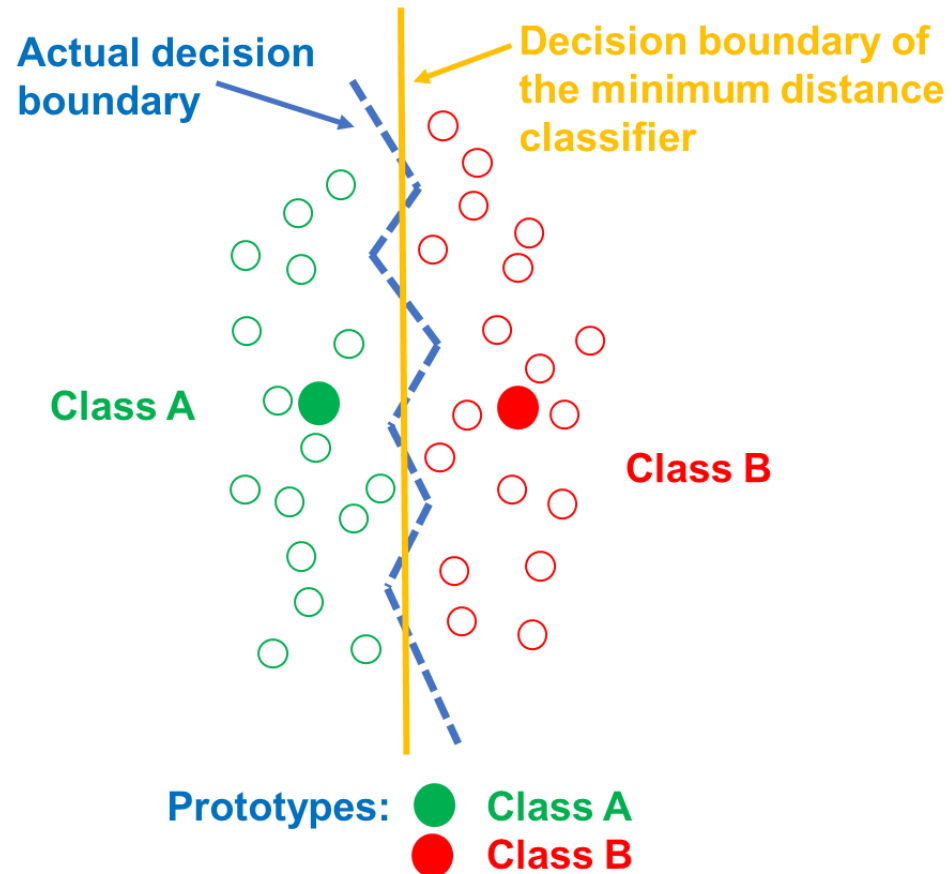
class(x) $\leftarrow \omega_i$



Minimum distance classifier (iii)

Drawback:

- Loss of useful information



k Nearest neighbours (k -NN)

- one of the **conceptually simplest** yet **powerful** models
- it takes **all the training data** into consideration
- it searches for the **k training data closest** to a test point
- a class label is assigned on the basis of a **majority vote**: the class that is most frequently represented among the k neighbours
- it is said to be a **lazy learner** since it does not perform any **training**. Instead, it just stores the data


k Nearest neighbours (k -NN) (ii)

We need to define two hyperparameters

- the **distance metric**, to determine which training data are the closest to a given test point
- the **k value**, which defines how many neighbours will be checked to guess the classification of a test point

k Nearest neighbours (k -NN) (iii)

Selecting the k value is the most critical question



k Nearest neighbours (k -NN) (iv)

Using error curves is a common method to choose the value of k

- Training error is zero at $k = 1$ because the nearest neighbour to a point is that point itself
- Test error is high at low values of k due to variance (overfitting), it subsequently lowers and stabilises, and with further increase in the value of k , the test error increases again due to bias (underfitting)



The value of k at which the **test error stabilises and is low** is taken as the **optimal value of k**

Similarity measures (i)

The distance metric:

- **Minkowski** distance: the **generalized form** of the Euclidean ($p = 2$), Manhattan ($p = 1$) and Chebyshev ($p = \infty$) distance metrics
 - **Euclidean (or L2 norm)** distance: the most commonly used distance measure, and it is limited to real-valued vectors
 - **Manhattan (or city block or L1 norm)** distance: measures the absolute value between two points
 - **Chebyshev (or chessboard)** distance: the minimum number of moves needed by a king to go from one square on a chessboard to another equals the Chebyshev distance

Similarity measures (ii)

The distance metric:

- **Hamming** distance: typically used with Boolean and string vectors
- **Cosine similarity**: defined as the cosine of the angle between two attribute vectors

Similarity measures (iii)

Minkowski distance:

$$d_p(x, y) = \left(\sum_{i=1}^d |y_i - x_i| \right)^{1/p}$$

Euclidean distance:

$$d_2(x, y) = \sqrt{\sum_{i=1}^d (y_i - x_i)^2}$$

Manhattan distance:


$$d_1(x, y) = \sum_{i=1}^d |y_i - x_i|$$

Chebyshev distance:

$$d_\infty(x, y) = \max_{i=1 \dots d} |y_i - x_i|$$

Similarity measures (iv)

An example of the Minkowski distances:



Manhattan distance = 12
(red, blue, or yellow)

Euclidean distance = 8.5
(green – continuous)

Chebyshev distance = 6
(green – discrete)

Similarity measures (v)

- **Hamming** distance:

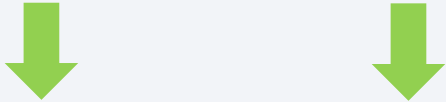
$$d_H(x, y) = \sum_{i=1}^d |y_i - x_i|$$

$$\begin{aligned} \text{if } x = y &\Rightarrow d_H(x, y) = 0 \\ \text{if } x \neq y &\Rightarrow d_H(x, y) = 1 \end{aligned}$$

Similarity measures (vi)

An example of the Hamming distance:

$$d_H(x, y) = \sum_{i=1}^d |y_i - x_i|$$



The diagram shows two vectors, Vector 1 and Vector 2, each with 8 elements. Two green arrows point down to the third and seventh elements of both vectors, which are highlighted in blue in the table below.

| | | | | | | | | |
|----------|---|---|---|---|---|---|---|---|
| Vector 1 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | 1 |
| Vector 2 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |

$$d_H = 0+0+1+0+0+0+1+0 = 2$$

the Hamming distance is 2
because **only two elements**
of the vectors differ

Similarity measures (vii)

- Cosine similarity:

$$S_C(\vec{x}, \vec{y}) = \cos(\theta) = \frac{\sum_{i=1}^d x_i y_i}{\sqrt{\sum_{i=1}^d x_i^2} \sqrt{\sum_{i=1}^d y_i^2}}$$

if $S_C(\vec{x}, \vec{y}) = -1 \Rightarrow \vec{x}$ and \vec{y} are exactly opposite

if $S_C(\vec{x}, \vec{y}) = +1 \Rightarrow \vec{x}$ and \vec{y} are exactly the same

if $S_C(\vec{x}, \vec{y}) = 0 \Rightarrow \vec{x}$ and \vec{y} are orthogonal

intermediate values indicate intermediate similarity

Normalization

- **Continuous numerical attributes:** to avoid that some attributes dominate over others, we should generally “normalize” the attributes in a common range such as $[0,1]$ or $[-1, 1]$
- **Normalization** is generally required when we are dealing with attributes on a **different scale** as this may lead to poor model performance

| person_name | Salary | Year_of_experience | Expected Position Level |
|-------------|--------|--------------------|-------------------------|
| Aman | 100000 | 10 | 2 |
| Abhinav | 78000 | 7 | 4 |
| Ashutosh | 32000 | 5 | 8 |
| Dishi | 55000 | 6 | 7 |
| Abhishek | 92000 | 8 | 3 |
| Avantika | 120000 | 15 | 1 |
| Ayushi | 65750 | 7 | 5 |

The attributes salary and year_of_experience are on different scale and hence attribute salary can take high priority over attribute year_of_experience in the model.

Normalization (ii): typical methods

- **z-score (or zero-mean) normalization or standardization:** it re-scales a feature value so that it has distribution with 0 mean value and variance equals to 1

$$z_i = \frac{x_i - \bar{x}_i}{\sigma_i}$$

where \bar{x}_i and σ_i are the mean value and the standard deviation of attribute i , respectively

- **min-max normalization:** it re-scales a feature or observation value with distribution value between 0 and 1

$$z_i = \frac{x_i - \min_i X}{\max_i X - \min_i X}$$

Normalization (iii): typical methods

- **decimal scaling:** it normalizes by moving the decimal point of values of the data. We divide each value of the data by the maximum absolute value of data

$$z_i = \frac{x_i}{10^j}$$

where j is the smallest integer such that $\max_i(|z_i|) < 1$

Example: our input data is -10, 201, 301, -401, 501, 601, 701

Step 1: Maximum absolute value in input data: 701

Step 2: Divide the input data by 1000 (i.e., $j = 3$, 10^3)


Result: The normalized data is: -0.01, 0.201, 0.301, -0.401, 0.501, 0.601, 0.701

The k -NN algorithm

1. Select the number k of neighbours
2. Calculate the distance between the training data and the test sample
3. Take the k nearest neighbours as per the calculated distance
4. Among these k neighbours, count the number of the training data in each class
5. Assign the test sample to that category (class) with a majority of neighbours

The k -NN algorithm (ii)

Example: k -NN classifier ($k = 5$ and **Euclidean distance**)



Advantages k -NN

Advantages:

- **Easy to understand and implement:** given its simplicity, it is one of the first classifiers that a new data scientist will learn
- **Adapts easily:** as new training samples are added, the algorithm adjusts to account for any new data since all training data are stored into memory
- **Few hyperparameters:** only requires a k value and a distance metric, which is low when compared to other machine learning models

Disadvantages k -NN

Disadvantages:

- **Does not scale well:** since k -NN is a lazy algorithm, it uses all the training data at the runtime and hence is slow
- **Curse of dimensionality:** the k -NN algorithm does not perform well with high-dimensional data
- **Complexity:** $O(n)$ for each instance to be classified
- **Computationally expensive:** it takes up more memory and data storage compared to other classifiers

The (k, l) -NN algorithm

Let l be a positive integer $\lceil k/2 \rceil < l \leq k$, a threshold for the majority in the voting of the k nearest neighbours

- This classifier consists of applying the k -NN model and make a decision based on:
 - if **some class has received at least l votes**, then assign the test sample to the most voted class
 - otherwise, reject the classification of the test sample (i.e., **assign it to a dummy class, ω_0**)
- This is a **classifier with reject option**

Application of k -NN: the IRIS data set

| Sepal Length | Sepal Width | Species | Distance | Rank |
|--------------|-------------|------------|----------|------|
| 5.3 | 3.7 | Setosa | | |
| 5.1 | 3.8 | Setosa | | |
| 7.2 | 3.0 | Virginica | | |
| 5.4 | 3.4 | Setosa | | |
| 5.1 | 3.3 | Setosa | | |
| 5.4 | 3.9 | Setosa | | |
| 7.4 | 2.8 | Virginica | | |
| 6.1 | 2.8 | Versicolor | | |
| 7.3 | 2.9 | Virginica | | |
| 6.0 | 2.7 | Versicolor | | |
| 5.8 | 2.8 | Virginica | | |
| 6.3 | 2.3 | Versicolor | | |
| 5.1 | 2.5 | Versicolor | | |
| 6.3 | 2.5 | Versicolor | | |

Application of k -NN: the IRIS data set (ii)

| Sepal Leght | Sepal Width | Species |
|-------------|-------------|---------|
| 5.2 | 3.1 | ? |

Compute the distances between the given test sample and all training data:

$$\begin{aligned}d_2(x, y) &= \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2} = \\ &= \sqrt{(5.2 - 5.3)^2 + (3.1 - 3.7)^2} = 0.608\end{aligned}$$

Application of k -NN: the IRIS data set (iii)

| Sepal Length | Sepal Width | Species | Distance | Rank |
|--------------|-------------|------------|----------|------|
| 5.3 | 3.7 | Setosa | 0.608 | |
| 5.1 | 3.8 | Setosa | 0.707 | |
| 7.2 | 3.0 | Virginica | 2.002 | |
| 5.4 | 3.4 | Setosa | 0.360 | |
| 5.1 | 3.3 | Setosa | 0.220 | |
| 5.4 | 3.9 | Setosa | 0.820 | |
| 7.4 | 2.8 | Virginica | 2.220 | |
| 6.1 | 2.8 | Versicolor | 0.940 | |
| 7.3 | 2.9 | Virginica | 2.100 | |
| 6.0 | 2.7 | Versicolor | 0.890 | |
| 5.8 | 2.8 | Virginica | 0.670 | |
| 6.3 | 2.3 | Versicolor | 1.360 | |
| 5.1 | 2.5 | Versicolor | 0.600 | |
| 6.3 | 2.5 | Versicolor | 1.250 | |

Application of *k*-NN: the IRIS data set (iv)

| Sepal Length | Sepal Width | Species | Distance | Rank |
|--------------|-------------|------------|----------|------|
| 5.3 | 3.7 | Setosa | 0.608 | 3 |
| 5.1 | 3.8 | Versicolor | 0.707 | 6 |
| 7.2 | 3.0 | Virginica | 2.002 | 12 |
| 5.4 | 3.4 | Setosa | 0.360 | 2 |
| 5.1 | 3.3 | Setosa | 0.220 | 1 |
| 5.4 | 3.9 | Setosa | 0.820 | 7 |
| 7.4 | 2.8 | Virginica | 2.220 | 14 |
| 6.1 | 2.8 | Versicolor | 0.940 | 9 |
| 7.3 | 2.9 | Virginica | 2.100 | 13 |
| 6.0 | 2.7 | Versicolor | 0.890 | 8 |
| 5.8 | 2.8 | Virginica | 0.670 | 5 |
| 6.3 | 2.3 | Versicolor | 1.360 | 11 |
| 5.1 | 2.5 | Versicolor | 0.600 | 4 |
| 6.3 | 2.5 | Versicolor | 1.250 | 10 |

Find ranks

Application of k -NN: the IRIS data set (v)

| Sepal Leght | Sepal Width | Species |
|-------------|-------------|---------|
| 5.2 | 3.1 | ? |

Find the k nearest neighbours:

| | | |
|-------------|----------------------|------------|
| If $k = 1$ | Rank = 1 | Setosa |
| If $k = 5$ | Rank = 1, 2, ..., 5 | Setosa |
| If $k = 10$ | Rank = 1, 2, ..., 10 | Versicolor |