



Multi-classifiers

Department of Computer Languages and Systems

Terminologies

- Lots of **terms** are used **to refer to multi-classifiers**:
 - ensemble of classifiers
 - combining classifiers
 - decision committee
 - multiple classifier system
 - mixture of experts
 - committee-based learning
 - etc.


Introduction: motivation

- When you have to face a complex classification problem:
 - which learning algorithm to use?
 - which parameters to choose?
 - how to use the training data?
 - which vector space to map the data onto? What is the most discriminating representation?

Introduction: motivation

- Different models may appear while searching for a solution, but often **none of them is better than the rest**
 - In this case, a reasonable choice is to keep them all and **create a final system integrating the pieces**
 - The core idea behind this is to aggregate multiple models to obtain **a combined model D that outperforms every single model D_i in it**
 - Each single model D_i is called **base learner** (classifier) or **individual learner** (classifier)

Strategies to build a multi-classifier



- Combination level: design **different combiners**
- Classifier level: use **different base classifiers**
- Data level: use **different data subsets**
- Feature level: use **different feature subsets**

Combination level: fusion vs. selection

Fusion

- each ensemble member is supposed to have knowledge of the whole feature space
- some combiner such as the average and majority vote is applied to label the input object x

Selection

- each ensemble member is supposed to know well a part of the feature space and to be responsible for objects in this part
- one member is chosen to label the input object x

Combination level (ii): fusion vs. selection

Fusion

- competitive classifiers
- ensemble approach
- multiple topology

Selection

- cooperative classifiers
- modular approach
- hybrid topology

Fusion: Majority vote

Decision rule: to choose the class most voted by the base classifiers

Three consensus patterns:

- **Unanimity** (all agree) 
- **Simple majority** (50%+1) 
- **Plurality** (most votes) 

Fusion: Majority vote (ii)

Let it be

- $[d_{i,1}, \dots, d_{i,C}]^T \in \{0,1\}^C, i = 1, \dots, L$, where $d_{i,j} = 1$ if D_i labels x in class ω_j , and 0 otherwise

Then, the **plurality vote rule** will result in an ensemble decision for class ω_k if

$$\sum_{i=1}^L d_{i,k} = \max_{j=1,\dots,C} \sum_{i=1}^L d_{i,j}$$

This rule coincides with the **simple majority rule** if $C = 2$

Fusion: Majority vote (iii)

A **thresholded plurality vote**: we increase the set of classes with one more class ω_{c+1} , for objects for which the ensemble does not determine a class label with a sufficient confidence. Now, the decision is

$$\begin{cases} \omega_k, & \text{if } \sum_{i=1}^L d_{i,k} \geq \alpha \cdot L \\ \omega_{c+1}, & \text{otherwise} \end{cases}$$

where $0 < \alpha \leq 1$. If $\alpha = 1$, this becomes the **unanimity vote rule**

Fusion: Majority vote (iv)

Weighted majority vote:

- an adequate option when the base classifiers are not of very similar accuracy
- it attempts to give the more competent classifiers more power in making the final decision

Fusion: Majority vote (v)

Weighted majority vote:

- we can represent the outputs as

$$d_{i,j} = \begin{cases} 1 & \text{if } D_i \text{ labels } x \text{ in } \omega_j \\ 0 & \text{otherwise} \end{cases}$$

- then, the decision is ω_k if

$$\sum_{i=1}^L w_i d_{i,k} = \max_{j=1,\dots,c} \sum_{i=1}^L w_i d_{i,j}$$

where $w_i \geq 0$ ($\sum_{i=1}^c w_i = 1$) is a weight for classifier D_i


Selection

Suppose an ensemble $D = \{D_1, \dots, D_L\}$ of classifiers already trained. Then, the feature space \mathbb{R}^d is divided into $K > 1$ **selection regions** (or **regions of competence**), which are denoted by R_1, \dots, R_K

- usually, $K = L$
- each region R_i is associated with a classifier, which will be responsible for deciding on the input objects in this part of the space
- these regions are not associated with specific classes, nor do they need to be of a certain shape or size

Selection (ii)

Example: suppose a data set with 2000 points and two classes ω_1 and ω_2 , and we have an ensemble with three classifiers D_1 , D_2 , D_3 , each one associated with regions R_1 , R_2 , R_3



- D_1 always predicts ω_1
- D_2 always predicts ω_2
- D_3 is a linear classifier whose discriminant function is shown as a dashed line
- Accuracy of the individual classifiers or that of a majority vote (fusion) is approximately 0.5
- Accuracy of the selection combiner will be close to 1

Classifier level: stacking

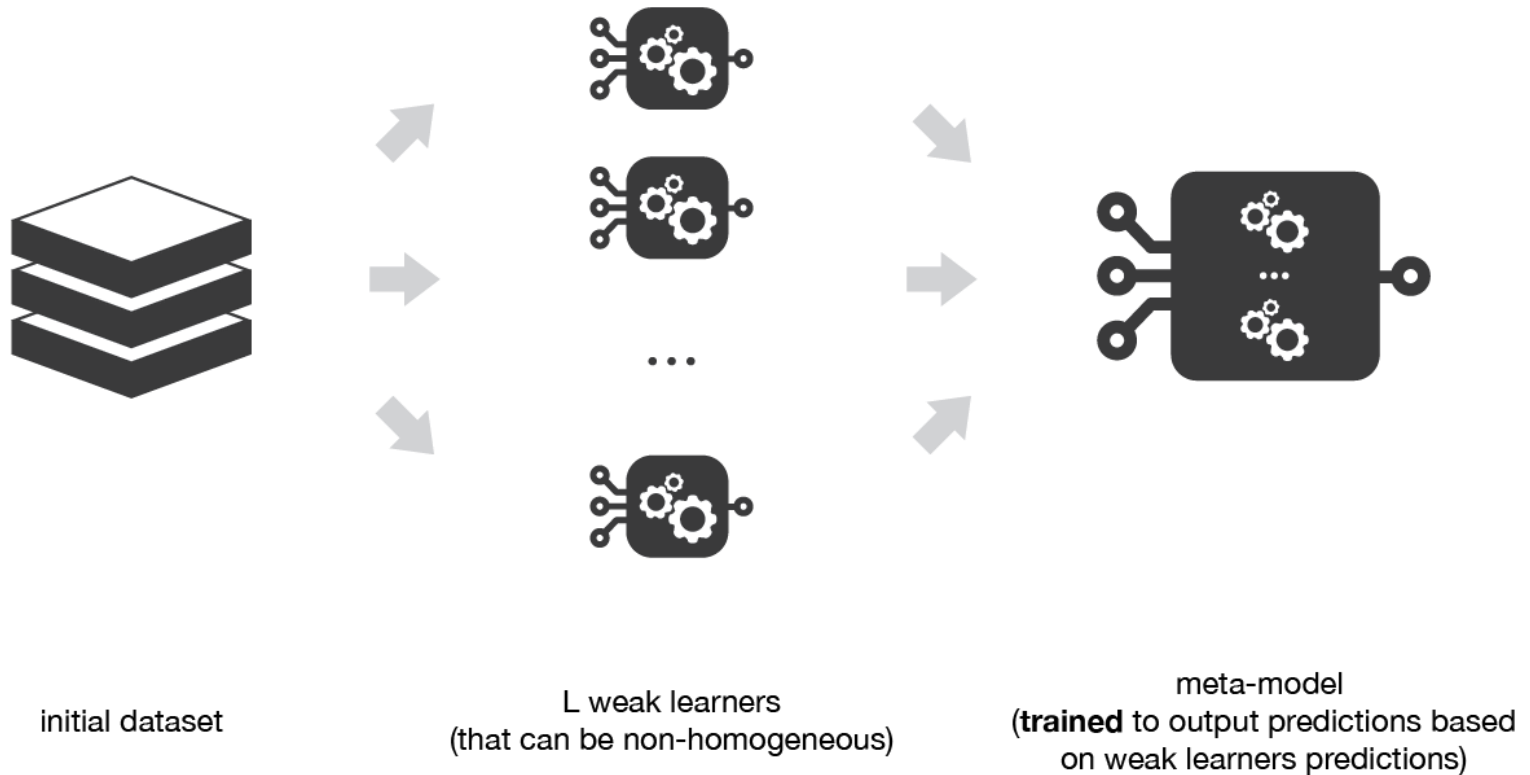
Idea:

- learn **various different weak learners** (base learners)
- combine the base learners by training a **meta-model**

Comments:

- we need to define two things in order to build our stacking model: **the L base learners we want to fit and the meta-model that combines them**
- for example, we can choose as weak learners a k -NN classifier, a decision tree and a SVM, and decide to learn a neural network as meta-model. Then, the neural network will take as inputs the outputs of our three weak learners and will learn to return final predictions based on it

Classifier level: stacking (ii)



Classifier level: stacking (iii)

1. Initialize the parameters
 L , the number of weak learners
2. Split the data into two folds
3. For $l = 1, \dots, L$
Train the weak learner to data of the first fold
Make predictions for data in the second fold
4. Train the meta-model on the second fold, using predictions made by the weak learners as inputs

Data level: bagging


Idea:

- the ensemble is made of classifiers built on **bootstrap replicates** of the training set $T_{tra} = \{x_1, \dots, x_n\}$
- the classifier outputs are combined by the **plurality vote**

Comments:

- we **sample with replacement** from the original T_{tra} to create L new training sets (often, also of size n)
- all L base classifiers are the **same classification model**
- the base classifier should be **unstable** (small changes in T_{tra} lead to large changes in the classifier output (**neural networks and decision trees are unstable, k -NN is stable**))
- this is a **parallel algorithm** in both its training and operational phases

Data level (ii): bagging



Data level (iii): bagging

Training phase

1. Initialize the parameters
 $D = \emptyset$, the ensemble
 L , the number of classifiers to train
2. For $l = 1, \dots, L$
Take a bootstrap sample S_l from the original training set T_{tra}
Build a classifier D_l using S_l as the training set
Add the classifier to the current ensemble, $D = D \cup D_l$
Return D

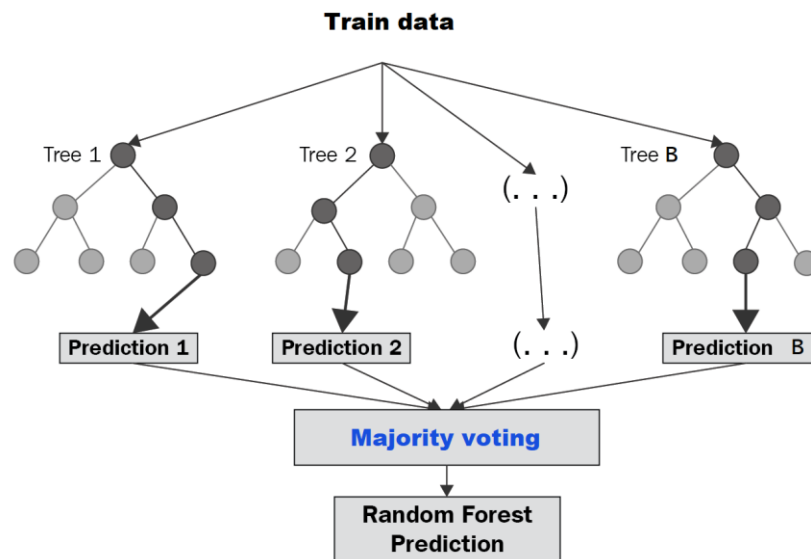
Classification (regression) phase

1. Run $D_1 \dots D_L$ on the input \mathbf{x}
2. Assign \mathbf{x} to the class with the maximum number of votes (simple majority voting, for classification)
Assign \mathbf{x} with the average of the estimated values (simple average, for regression)

Data level (iv): variants of bagging

Random forest

- a collection of full decision trees built in parallel from random bootstrap sample of the data set
- the final prediction is an average of all of the decision tree predictions



Data level (vi): boosting

Idea:


- to develop the ensemble D incrementally, adding one base classifier at a time
- some classifiers have more say in the classification than others
- the classifier D_i is made by taking the errors of the classifier D_{i-1} into account

Comments:

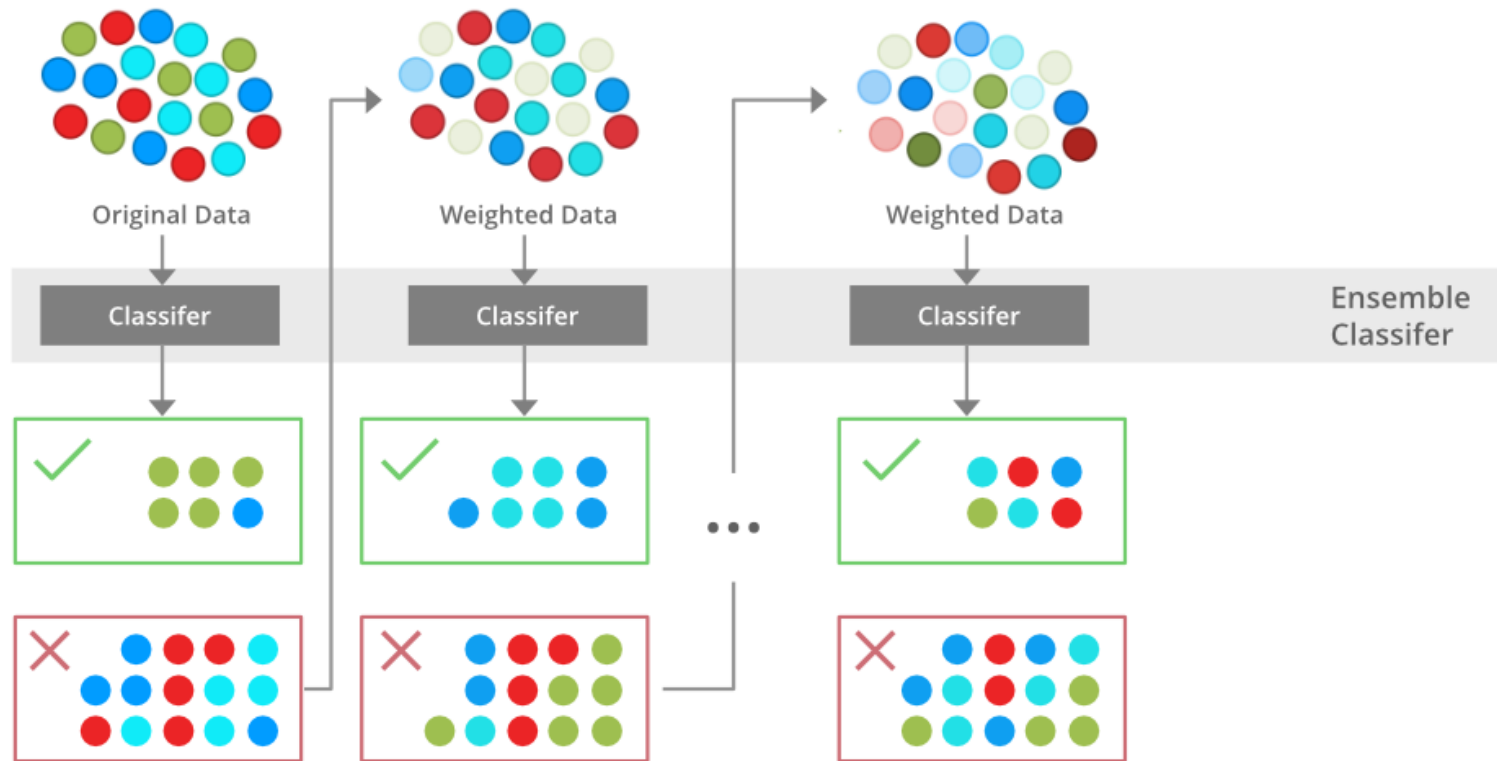
- this is a sequential algorithm
- the errors that the first classifier makes influence how the second classifier is made, and so on

Data level (vii): boosting

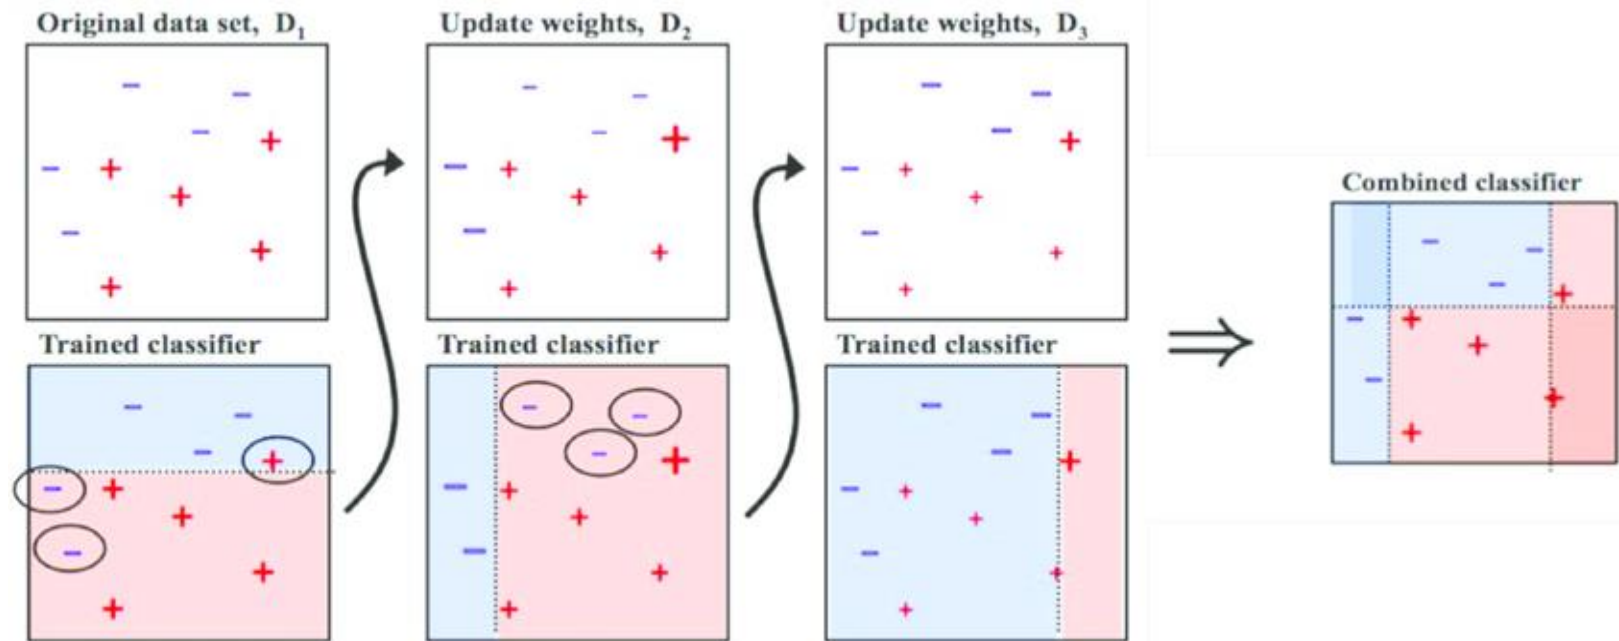
The idea of boosting could be seen as a golfer who initially hits a golf ball towards the hole at position y , but only goes as far as f_0 . The golfer then repeatedly hits the ball more gently, moving it toward the hole a little at a time and after reassessing the direction and distance to the hole with each shot.



Data level (viii): boosting



Data level (ix): boosting



Data level (x): boosting (AdaBoost)

Training phase

1. Initialize the parameters

Set the weights $w^i = 1/n$ (equal weights to each data point)

$D = \emptyset$, the ensemble

L , the number of classifiers

2. For $l = 1, \dots, L$

Build a classifier D_l with the training data using w^i for $i = 1, \dots, n$

Calculate the proportion of errors in classification e_l

Compute $S_l = \log((1 - e_l)/e_l)$

Update the weights w^i (weights of correctly classified samples do not change; incorrectly classified samples are given more weight by multiplying their previous weight by $(1 - e_l)/e_l$)

Data level (xi): boosting (AdaBoost)

Classification phase

Given a sample \mathbf{x} , if we denote $\hat{y}_l(\mathbf{x})$ its classification using classifier D_l , then

$$\hat{y}(\mathbf{x}) = \text{sign} \left(\sum_l S_l \hat{y}_l(\mathbf{x}) \right)$$

(if the sum is positive, the observation is classified as belonging to class +1, otherwise to class -1)

Data level (xii): variants of boosting

Gradient boosting

- it involves three elements:
 - a **loss function** to be optimized (e.g., regression may use mean squared error and classification may use logarithmic loss)
 - a **weak learner** to make predictions (usually, decision trees)
 - an additive model that minimizes the loss function when adding trees (**gradient descent** is used to minimize the loss)

Data level (xiii): variants of boosting

Extreme gradient boosting (XGBoost)

- an **efficient and effective** implementation of gradient boosting
- it is highly **scalable** and can handle **large data sets**
- **trees are built in parallel**, instead of sequentially like gradient boosting
- it implements **early stopping** so we can stop model evaluation when additional trees offer no improvement

Data level (xiv): variants of boosting

Categorical boosting (CatBoost)

- it is designed to **work on heterogeneous data** (categorical, numerical, logical, ...)
- it works well with **less data**
- improved accuracy by **reducing overfitting**

Feature level: random subspace

Idea:

- the ensemble is made of classifiers built on **random subsets of features (with replacement)** of predefined size d_{rs} ($d_{rs} < d$)
- the classifier outputs are combined by the **plurality vote**

Comments:

- an attractive choice for **high-dimensional problems** where the number of features (d) is much larger than the number of training points (n)
- it works best when the **discriminative information** is “**dispersed**” across all the features

Feature level (ii): random subspace

Training phase

1. Initialize the parameters
 $D = \emptyset$, the ensemble
 L , the number of classifiers to train
2. For $k = 1, \dots, L$
Pick up d_{rs} features from d with replacement
Build a classifier D_k using the subspace sample
Add the classifier to the current ensemble, $D = D \cup D_k$
3. Return D

Classification phase

1. Run D_1, \dots, D_L on the input x
2. Assign x to the class with the maximum number of votes