



# Multi-classifiers

Department of Computer Languages and Systems

# Terminologies

- Lots of **terms** are used **to refer to multi-classifiers**:
  - ensemble of classifiers
  - combining classifiers
  - decision committee
  - multiple classifier system
  - mixture of experts
  - committee-based learning
  - etc.


# Introduction: motivation

- When you have to face a complex classification problem:
  - which **learning algorithm** to use?
  - which **parameters** to choose?
  - how to **use the training data**?
  - which **vector space to map the data onto**? What is the **most discriminating representation**?

# Introduction: motivation

- Different models may appear while searching for a solution, but often none of them is better than the rest
  - In this case, a reasonable choice is to keep them all and create a final system integrating the pieces
  - The core idea behind this is to aggregate multiple models to obtain a combined model  $D$  that outperforms every single model  $D_i$  in it
  - Each single model  $D_i$  is called base learner (classifier) or individual learner (classifier)

# Strategies to build a multi-classifier



- Combination level: design **different combiners**
- Classifier level: use **different base classifiers**
- Data level: use **different data subsets**
- Feature level: use **different feature subsets**

# Combination level: fusion vs. selection

## Fusion

- each ensemble member is supposed to have knowledge of the whole feature space
- some combiner such as the average and majority vote is applied to label the input object  $x$

## Selection

- each ensemble member is supposed to know well a part of the feature space and to be responsible for objects in this part
- one member is chosen to label the input object  $x$

# Combination level (ii): fusion vs. selection

## Fusion

- competitive classifiers
- ensemble approach
- multiple topology

## Selection

- cooperative classifiers
- modular approach
- hybrid topology

# Fusion: Majority vote

**Decision rule:** to choose the class most voted by the base classifiers

Three consensus patterns:

- **Unanimity** (all agree) 
- **Simple majority** (50%+1) 
- **Plurality** (most votes) 



# Fusion: Majority vote (ii)

Let it be

- $[d_{i,1}, \dots, d_{i,C}]^T \in \{0,1\}^C, i = 1, \dots, L$ , where  $d_{i,j} = 1$  if  $D_i$  labels  $x$  in class  $\omega_j$ , and 0 otherwise

Then, the **plurality vote rule** will result in an ensemble decision for class  $\omega_k$  if

$$\sum_{i=1}^L d_{i,k} = \max_{j=1,\dots,C} \sum_{i=1}^L d_{i,j}$$

This rule coincides with the **simple majority rule** if  $C = 2$

## Fusion: Majority vote (iii)

A **thresholded plurality vote**: we increase the set of classes with one more class  $\omega_{c+1}$ , for objects for which the ensemble does not determine a class label with a sufficient confidence. Now, the decision is

$$\begin{cases} \omega_k, & \text{if } \sum_{i=1}^L d_{i,k} \geq \alpha \cdot L \\ \omega_{c+1}, & \text{otherwise} \end{cases}$$

where  $0 < \alpha \leq 1$ . If  $\alpha = 1$ , this becomes the **unanimity vote rule**

# Fusion: Majority vote (iv)

## Weighted majority vote:

- an adequate option when the base classifiers are not of very similar accuracy
- it attempts to give the more competent classifiers more power in making the final decision

# Fusion: Majority vote (v)

## Weighted majority vote:

- we can represent the outputs as

$$d_{i,j} = \begin{cases} 1 & \text{if } D_i \text{ labels } x \text{ in } \omega_j \\ 0 & \text{otherwise} \end{cases}$$

- then, the decision is  $\omega_k$  if

$$\sum_{i=1}^L w_i d_{i,k} = \max_{j=1,\dots,c} \sum_{i=1}^L w_i d_{i,j}$$

where  $w_i \geq 0$  ( $\sum_{i=1}^c w_i = 1$ ) is a weight for classifier  $D_i$


# Selection

Suppose an ensemble  $D = \{D_1, \dots, D_L\}$  of classifiers already trained. Then, the feature space  $\mathbb{R}^d$  is divided into  $K > 1$  **selection regions** (or **regions of competence**), which are denoted by  $R_1, \dots, R_K$

- usually,  $K = L$
- each region  $R_i$  is associated with a classifier, which will be responsible for deciding on the input objects in this part of the space
- these regions are not associated with specific classes, nor do they need to be of a certain shape or size

# Selection (ii)

**Example:** suppose a data set with 2000 points and two classes  $\omega_1$  and  $\omega_2$ , and we have an ensemble with three classifiers  $D_1$ ,  $D_2$ ,  $D_3$ , each one associated with regions  $R_1$ ,  $R_2$ ,  $R_3$



- $D_1$  always predicts  $\omega_1$
- $D_2$  always predicts  $\omega_2$
- $D_3$  is a linear classifier whose discriminant function is shown as a dashed line
- Accuracy of the individual classifiers or that of a majority vote (fusion) is approximately 0.5
- Accuracy of the selection combiner will be close to 1

# Classifier level: stacking

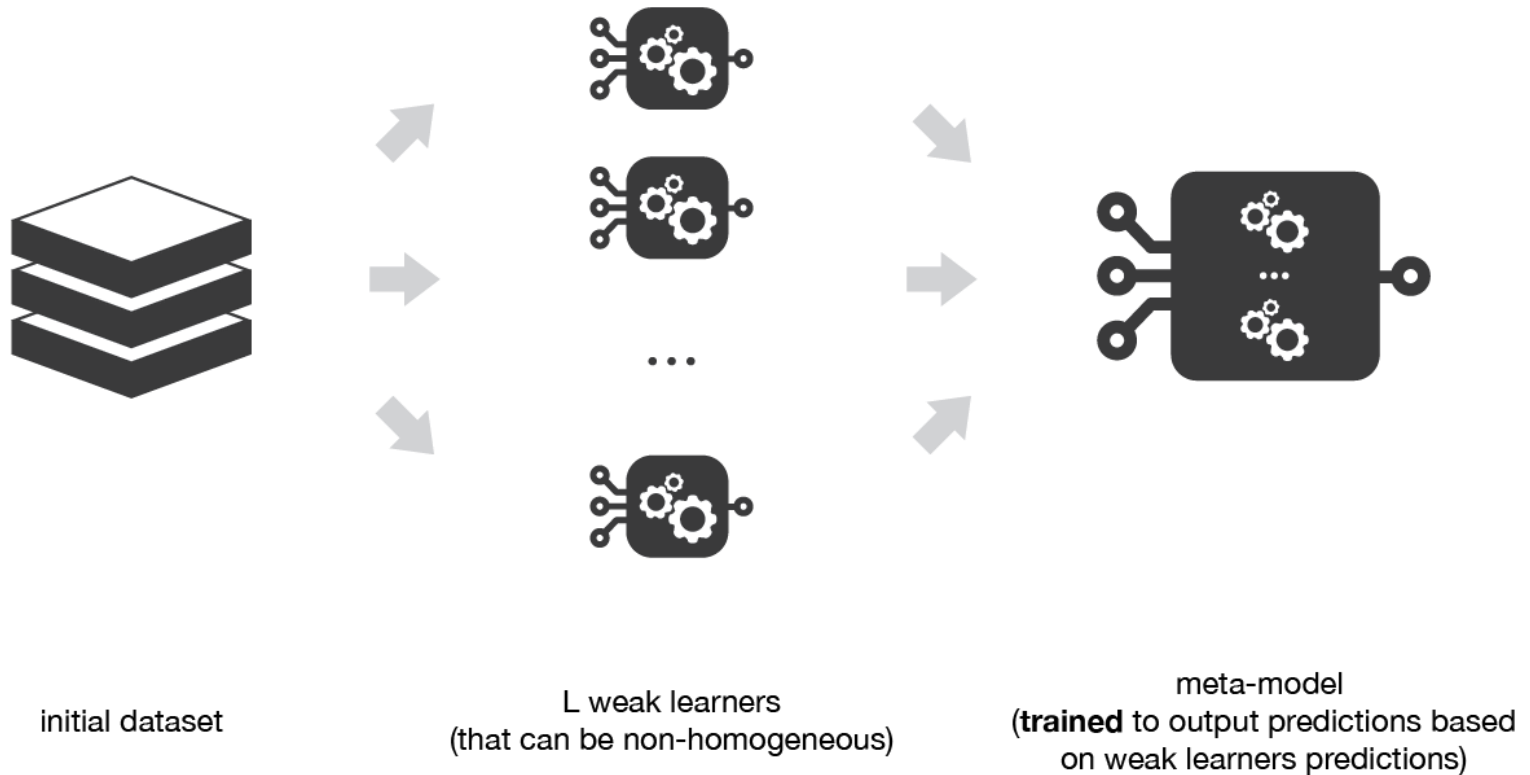
## Idea:

- learn **various different weak learners** (base learners)
- combine the base learners by training a **meta-model**

## Comments:

- we need to define two things in order to build our stacking model: **the  $L$  base learners we want to fit and the meta-model that combines them**
- for example, we can choose as weak learners a  $k$ -NN classifier, a decision tree and a SVM, and decide to learn a neural network as meta-model. Then, the neural network will take as inputs the outputs of our three weak learners and will learn to return final predictions based on it

# Classifier level: stacking (ii)





# Classifier level: stacking (iii)

1. Initialize the parameters  
 $L$ , the number of weak learners
2. Split the data into two folds
3. For  $l = 1, \dots, L$   
Train the weak learner to data of the first fold  
Make predictions for data in the second fold
4. Train the meta-model on the second fold, using predictions made by the weak learners as inputs

# Data level: bagging


## Idea:

- the ensemble is made of classifiers built on **bootstrap replicates** of the training set  $T_{tra} = \{x_1, \dots, x_n\}$
- the classifier outputs are combined by the **plurality vote**

## Comments:

- we **sample with replacement** from the original  $T_{tra}$  to create  $L$  new training sets (often, also of size  $n$ )
- all  $L$  base classifiers are the **same classification model**
- the base classifier should be **unstable** (small changes in  $T_{tra}$  lead to large changes in the classifier output (**neural networks and decision trees are unstable,  $k$ -NN is stable**))
- this is a **parallel algorithm** in both its training and operational phases

# Data level (ii): bagging



# Data level (iii): bagging

## Training phase

1. Initialize the parameters  
 $D = \emptyset$ , the ensemble  
 $L$ , the number of classifiers to train
2. For  $l = 1, \dots, L$   
Take a bootstrap sample  $S_l$  from the original training set  $T_{tra}$   
Build a classifier  $D_l$  using  $S_l$  as the training set  
Add the classifier to the current ensemble,  $D = D \cup D_l$   
Return  $D$

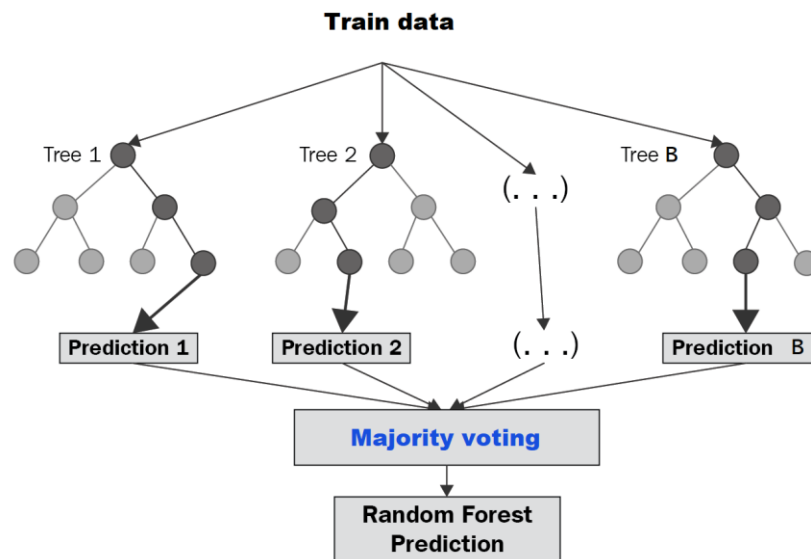
## Classification (regression) phase

1. Run  $D_1 \dots D_L$  on the input  $\mathbf{x}$
2. Assign  $\mathbf{x}$  to the class with the maximum number of votes (simple majority voting, for classification)  
Assign  $\mathbf{x}$  with the average of the estimated values (simple average, for regression)

# Data level (iv): variants of bagging

## Random forest

- a collection of full decision trees built in parallel from random bootstrap sample of the data set
- the final prediction is an average of all of the decision tree predictions



# Data level (vi): boosting

## Idea:


- to develop the ensemble  $D$  incrementally, adding one base classifier at a time
- some classifiers have more say in the classification than others
- the classifier  $D_i$  is made by taking the errors of the classifier  $D_{i-1}$  into account

## Comments:

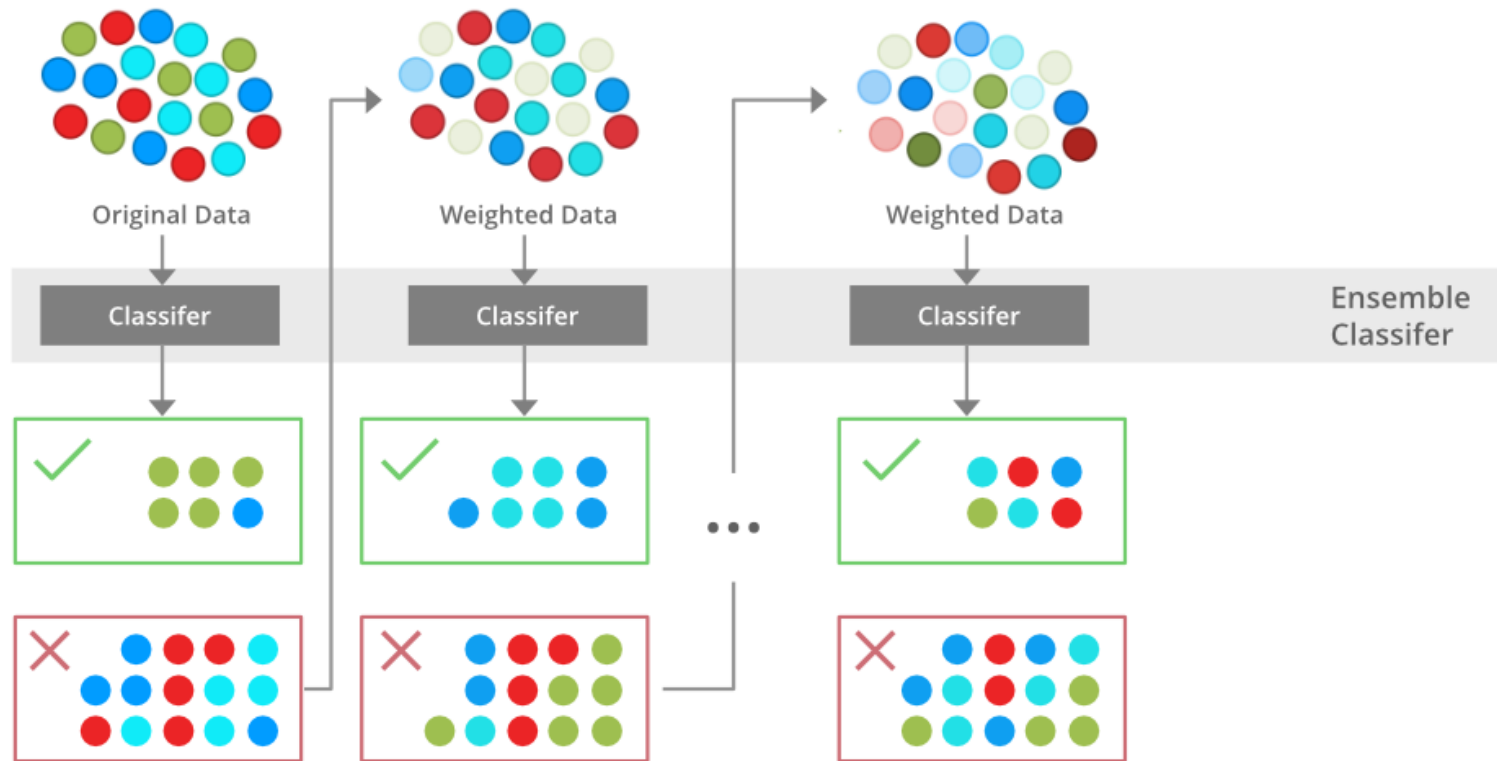
- this is a sequential algorithm
- the errors that the first classifier makes influence how the second classifier is made, and so on

# Data level (vii): boosting

The idea of boosting could be seen as a golfer who initially hits a golf ball towards the hole at position  $y$ , but only goes as far as  $f_0$ . The golfer then repeatedly hits the ball more gently, moving it toward the hole a little at a time and after reassessing the direction and distance to the hole with each shot.

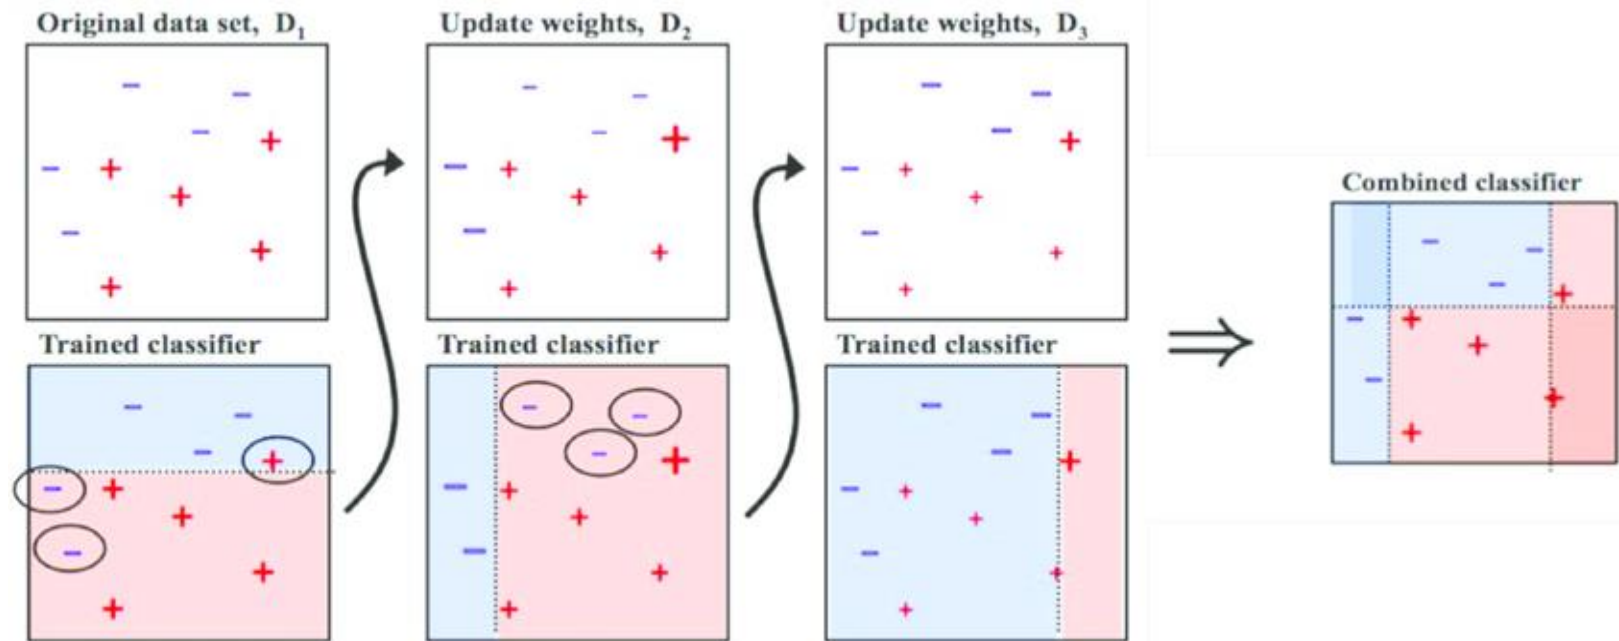


# Data level (viii): boosting





# Data level (ix): boosting



# Data level (x): boosting (AdaBoost)

## Training phase

1. Initialize the parameters

Set the weights  $w^i = 1/n$  (equal weights to each data point)

$D = \emptyset$ , the ensemble

$L$ , the number of classifiers

2. For  $l = 1, \dots, L$

Build a classifier  $D_l$  with the training data using  $w^i$  for  $i = 1, \dots, n$

Calculate the proportion of errors in classification  $e_l$

Compute  $S_l = \log((1 - e_l)/e_l)$

Update the weights  $w^i$  (weights of correctly classified samples do not change; incorrectly classified samples are given more weight by multiplying their previous weight by  $(1 - e_l)/e_l$ )

# Data level (xi): boosting (AdaBoost)

## Classification phase

Given a sample  $\mathbf{x}$ , if we denote  $\hat{y}_l(\mathbf{x})$  its classification using classifier  $D_l$ , then

$$\hat{y}(\mathbf{x}) = \text{sign} \left( \sum_l S_l \hat{y}_l(\mathbf{x}) \right)$$

(if the sum is positive, the observation is classified as belonging to class +1, otherwise to class -1)

# Data level (xii): variants of boosting

## Gradient boosting

- it involves three elements:
  - a **loss function** to be optimized (e.g., regression may use mean squared error and classification may use logarithmic loss)
  - a **weak learner** to make predictions (usually, decision tress)
  - an additive model that minimizes the loss function when adding trees (**gradient descent** is used to minimize the loss)

# Data level (xiii): variants of boosting

## Extreme gradient boosting (XGBoost)

- an **efficient and effective** implementation of gradient boosting
- it is highly **scalable** and can handle **large data sets**
- **trees are built in parallel**, instead of sequentially like gradient boosting
- it implements **early stopping** so we can stop model evaluation when additional trees offer no improvement

# Data level (xiv): variants of boosting

## Categorical boosting (CatBoost)

- it is designed to **work on heterogeneous data** (categorical, numerical, logical, ...)
- it works well with **less data**
- improved accuracy by **reducing overfitting**

# Feature level: random subspace

## Idea:

- the ensemble is made of classifiers built on **random subsets of features (with replacement)** of predefined size  $d_{rs}$  ( $d_{rs} < d$ )
- the classifier outputs are combined by the **plurality vote**

## Comments:

- an attractive choice for **high-dimensional problems** where the number of features ( $d$ ) is much larger than the number of training points ( $n$ )
- it works best when the **discriminative information** is “**dispersed**” across all the features

# Feature level (ii): random subspace

## Training phase

1. Initialize the parameters  
 $D = \emptyset$ , the ensemble  
 $L$ , the number of classifiers to train
2. For  $k = 1, \dots, L$   
Pick up  $d_{rs}$  features from  $d$  with replacement  
Build a classifier  $D_k$  using the subspace sample  
Add the classifier to the current ensemble,  $D = D \cup D_k$
3. Return  $D$

## Classification phase

1. Run  $D_1, \dots, D_L$  on the input  $x$
2. Assign  $x$  to the class with the maximum number of votes